

Course 2 Section 1.9 - Joins

Jiaying Wu

```
library(tidyverse)
```

Bank data: left join

Below are two made-up data sets: customer and transaction.

- customer contains some personal information about a bank's customers.
- transaction contains spending transactions from some customers of the bank

```
# Create bank customer info data
customer <- tibble(customer_id = c("C120463", "N4244552", "A907892", "Z337572", "D701453", "A285883"),
  names = c("John", "Sam", "Mike", "Rachael", "Mary", "Will"),
  sex = c("M", "F", "M", "F", "F", "M"),
  educ = c("HS", "Bachelor", "PhD", "Bachelor", "HS", "HS"),
  wage = c(1200, 900, 3100, 600, 500, 800))

# Print customer
customer
```

```
## # A tibble: 6 x 5
##   customer_id names    sex    educ    wage
##   <chr>        <chr>  <chr> <chr>    <dbl>
## 1 C120463     John    M     HS      1200
## 2 N4244552    Sam     F     Bachelor 900
## 3 A907892     Mike    M     PhD      3100
## 4 Z337572     Rachael F     Bachelor 600
## 5 D701453     Mary    F     HS        500
## 6 A285883     Will    M     HS        800
```

```
# Create bank spending transaction data
transaction <- tibble(customer_id = c("D701453", "N4244552", "C120463", "A907892", "D701453", "C120463"),
  shop = c("JB Hifi", "Steakhouse", "Apple", "Coles", "Lobster Diner", "Dymocks", "JB Hifi"),
  amount = c(300, 110, 3000, 80, 185, 40, 25, 15, 170))

# Print transaction
transaction
```

```
## # A tibble: 9 x 3
##   customer_id shop          amount
##   <chr>        <chr>        <dbl>
## 1 D701453     JB Hifi          300
```

```
## 2 N4244552    Steakhouse      110
## 3 C120463     Apple          3000
## 4 A907892     Coles             80
## 5 D701453     Lobster Diner    185
## 6 C120463     Dymocks          40
## 7 N4244552    Target            25
## 8 C120463     Netflix          15
## 9 A907892     Mecca            170
```

Q1: left join customer into transaction by the customer_id variable.

```
# Left join customer into transaction by customer_id
transaction_customer <- left_join(transaction, customer, by = "customer_id")

# Print transaction_customer
transaction_customer
```

```
## # A tibble: 9 x 7
##   customer_id shop          amount names sex   educ      wage
##   <chr>        <chr>        <dbl> <chr> <chr> <chr>    <dbl>
## 1 D701453     JB Hifi           300 Mary  F     HS        500
## 2 N4244552    Steakhouse       110 Sam   F     Bachelor  900
## 3 C120463     Apple           3000 John  M     HS       1200
## 4 A907892     Coles             80 Mike  M     PhD       3100
## 5 D701453     Lobster Diner   185 Mary  F     HS        500
## 6 C120463     Dymocks          40 John  M     HS       1200
## 7 N4244552    Target            25 Sam   F     Bachelor  900
## 8 C120463     Netflix          15 John  M     HS       1200
## 9 A907892     Mecca            170 Mike  M     PhD       3100
```

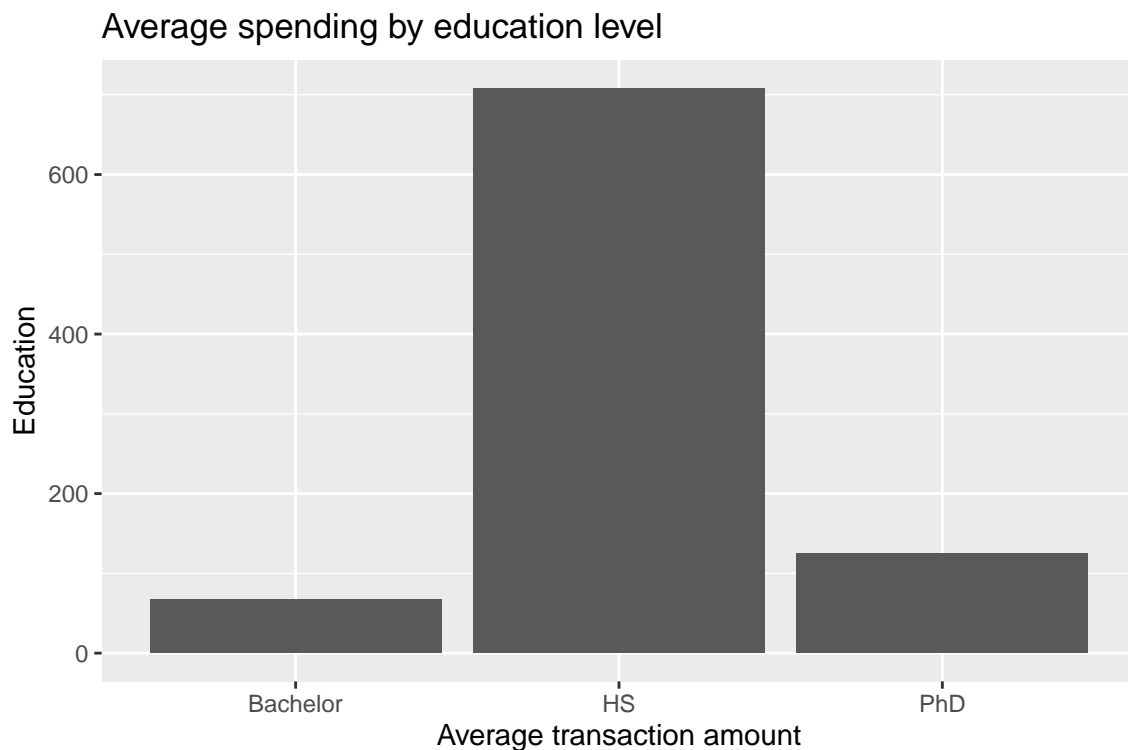
```
# Above join can also be done with the pipe operator (%>%)
transaction %>% left_join(customer, by = "customer_id")
```

```
## # A tibble: 9 x 7
##   customer_id shop          amount names sex   educ      wage
##   <chr>        <chr>        <dbl> <chr> <chr> <chr>    <dbl>
## 1 D701453     JB Hifi           300 Mary  F     HS        500
## 2 N4244552    Steakhouse       110 Sam   F     Bachelor  900
## 3 C120463     Apple           3000 John  M     HS       1200
## 4 A907892     Coles             80 Mike  M     PhD       3100
## 5 D701453     Lobster Diner   185 Mary  F     HS        500
## 6 C120463     Dymocks          40 John  M     HS       1200
## 7 N4244552    Target            25 Sam   F     Bachelor  900
## 8 C120463     Netflix          15 John  M     HS       1200
## 9 A907892     Mecca            170 Mike  M     PhD       3100
```

customer_id is the variable that connects both data sets together. Left joining customer into transaction using customer_id as the connecting variable add the variables names, sex, educ and wage into the transactions data.

Q2: One insight we can gather from this joined data set, which we have named `transaction_customer`, is the average amount spent by individuals of various education levels:

```
# Bar plot of average spending by education
transaction_customer %>%
  # group data by education
  group_by(educ) %>%
  # compute average spending (by education)
  summarise(avg_amount = mean(amount)) %>%
  ungroup() %>%
  # aesthetic for bar plot
  ggplot(aes(x = educ, y = avg_amount)) +
  # bar is the visual element we want
  # stat = "identity" to tell R to use y aesthetic and not it's own count for the y aesthetic
  geom_bar(stat = "identity") +
  # label the graph
  labs(title = "Average spending by education level",
       y = "Education",
       x = "Average transaction amount")
```



Joining the bank's customers' personal information to data about spending transactions is different from joining data about transactions to the personal details of banking customers.

Put differently, a left join of customer into transaction produces a joined data set that differs from a left join of transaction into customer.

Q3: perform a left join of transaction into customer

```
# Left join transaction into customer by customer_id
customer %>%
  left_join(transaction, by = "customer_id")
```

```
## # A tibble: 11 x 7
##   customer_id names    sex  educ    wage shop    amount
##   <chr>      <chr> <chr> <chr>   <dbl> <chr>    <dbl>
## 1 C120463    John    M    HS     1200 Apple     3000
## 2 C120463    John    M    HS     1200 Dymocks     40
## 3 C120463    John    M    HS     1200 Netflix     15
## 4 N4244552    Sam     F    Bachelor 900 Steakhouse 110
## 5 N4244552    Sam     F    Bachelor 900 Target     25
## 6 A907892    Mike    M    PhD     3100 Coles      80
## 7 A907892    Mike    M    PhD     3100 Mecca     170
## 8 Z337572    Rachael F    Bachelor 600 <NA>      NA
## 9 D701453    Mary    F    HS      500 JB Hifi    300
## 10 D701453    Mary    F    HS      500 Lobster Diner 185
## 11 A285883    Will    M    HS      800 <NA>      NA
```

Bank data: inner join

An inner join of customer and transaction keeps only observations with customer IDs that are in both the customer and transaction data.

Q1:

```
# Inner join customer and transaction by customer_id
customer %>%
  inner_join(transaction, by = "customer_id")
```

```
## # A tibble: 9 x 7
##   customer_id names    sex  educ    wage shop    amount
##   <chr>      <chr> <chr> <chr>   <dbl> <chr>    <dbl>
## 1 C120463    John    M    HS     1200 Apple     3000
## 2 C120463    John    M    HS     1200 Dymocks     40
## 3 C120463    John    M    HS     1200 Netflix     15
## 4 N4244552    Sam     F    Bachelor 900 Steakhouse 110
## 5 N4244552    Sam     F    Bachelor 900 Target     25
## 6 A907892    Mike    M    PhD     3100 Coles      80
## 7 A907892    Mike    M    PhD     3100 Mecca     170
## 8 D701453    Mary    F    HS      500 JB Hifi    300
## 9 D701453    Mary    F    HS      500 Lobster Diner 185
```

For inner joins, it does not matter which data set comes first, i.e, whether customer or transaction is placed before %>% does not produce an informatively different inner joined data. The only difference that we would see is how the columns and rows are arranged.

Notice that the first left join data is the same as this inner join data. Execute the following code chunk to inner join transaction and FAANG. The resulting inner join data is shown below.

```
# Create data of 2019 FAANG revenue
FAANG <- tibble(company = c("Facebook", "Amazon", "Apple", "Netflix", "Alphabet (formerly Google)"),
                revenue_billion = c(70.7, 280.5, 260.2, 20.16, 161.86))

# Print FAANG
FAANG
```

```
## # A tibble: 5 x 2
##   company                revenue_billion
##   <chr>                  <dbl>
## 1 Facebook                70.7
## 2 Amazon                 280.
## 3 Apple                  260.
## 4 Netflix                 20.2
## 5 Alphabet (formerly Google) 162.
```

Q2:

```
# Inner join transaction and FAANG by shop = company
transaction %>%
  inner_join(FAANG, by = c("shop" = "company"))
```

```
## # A tibble: 2 x 4
##   customer_id shop    amount revenue_billion
##   <chr>      <chr>    <dbl>      <dbl>
## 1 C120463    Apple      3000        260.
## 2 C120463    Netflix     15         20.2
```

Joining data from nycflights13

Previously, we explored the flights data set from the nycflights13 package. Other data sets included in the nycflights13 package include airlines and airports. Execute the following code chunk to explore all 3 data sets.

```
# Load nycflights13
library(nycflights13)
```

```
# Head of flights
head(flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>     <int>         <int>
## 1  2013     1     1     517           515             2       830           819
## 2  2013     1     1     533           529             4       850           830
## 3  2013     1     1     542           540             2       923           850
## 4  2013     1     1     544           545            -1      1004          1022
## 5  2013     1     1     554           600            -6       812           837
## 6  2013     1     1     554           558            -4       740           728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
# Head of airlines
head(airlines)
```

```
## # A tibble: 6 x 2
##   carrier name
##   <chr>   <chr>
## 1 9E      Endeavor Air Inc.
## 2 AA      American Airlines Inc.
## 3 AS      Alaska Airlines Inc.
## 4 B6      JetBlue Airways
## 5 DL      Delta Air Lines Inc.
## 6 EV      ExpressJet Airlines Inc.
```

```
# Head of airports
head(airports)
```

```
## # A tibble: 6 x 8
##   faa   name                lat   lon   alt   tz dst  tzone
##   <chr> <chr>                <dbl> <dbl> <dbl> <dbl> <chr> <chr>
## 1 04G   Lansdowne Airport      41.1 -80.6 1044   -5 A   America/New_Y~
## 2 06A   Moton Field Municipal Airp~ 32.5 -85.7 264    -6 A   America/Chica~
## 3 06C   Schaumburg Regional     42.0 -88.1 801    -6 A   America/Chica~
## 4 06N   Randall Airport        41.4 -74.4 523    -5 A   America/New_Y~
## 5 09J   Jekyll Island Airport   31.1 -81.4 11     -5 A   America/New_Y~
## 6 0A9   Elizabethton Municipal Air~ 36.4 -82.2 1593   -5 A   America/New_Y~
```

To add the full airline name to the flights data set, you can combine airlines to flights with the `left_join()` function, as shown in the following code chunk. Execute the following code chunks to left join airlines to flights and airports to flights.

```
# Left join airlines into flights
flights %>% left_join(airlines, by = "carrier")
```

```
## # A tibble: 336,776 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1 2013     1     1     517           515           2       830           819
## 2 2013     1     1     533           529           4       850           830
## 3 2013     1     1     542           540           2       923           850
## 4 2013     1     1     544           545          -1      1004          1022
## 5 2013     1     1     554           600          -6       812           837
## 6 2013     1     1     554           558          -4       740           728
## 7 2013     1     1     555           600          -5       913           854
## 8 2013     1     1     557           600          -3       709           723
## 9 2013     1     1     557           600          -3       838           846
## 10 2013     1     1     558           600          -2       753           745
## # ... with 336,766 more rows, and 12 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   name <chr>
```

```
# Left join airports into flights
flights %>% left_join(airports, by = c("origin" = "faa"))
```

```
## # A tibble: 336,776 x 26
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     517             515           2     830             819
## 2  2013     1     1     533             529           4     850             830
## 3  2013     1     1     542             540           2     923             850
## 4  2013     1     1     544             545          -1    1004            1022
## 5  2013     1     1     554             600          -6     812             837
## 6  2013     1     1     554             558          -4     740             728
## 7  2013     1     1     555             600          -5     913             854
## 8  2013     1     1     557             600          -3     709             723
## 9  2013     1     1     557             600          -3     838             846
## 10 2013     1     1     558             600          -2     753             745
## # ... with 336,766 more rows, and 18 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,
## #   name <chr>, lat <dbl>, lon <dbl>, alt <dbl>, tz <dbl>, dst <chr>,
## #   tzone <chr>
```